

Observe, Expand, Recover: An Observability-Guided Agent for Safe Query Expansion

Author One*, Author Two†

*Affiliation One, email@example.com

†Affiliation Two (AI Observability), email2@example.com

Abstract—Large language model (LLM) query expansion is typically applied either unconditionally or not at all, despite well-documented evidence that unconditional expansion is net-harmful on a non-trivial fraction of queries. We formulate query expansion instead as a bounded, observable agentic decision, evaluated under a strictly separated validation/test protocol: every retrieval-config, verifier-model, and threshold decision is fixed on a held-out TripClick TAIL *validation* split before a single frozen touch of *test*, per research question. A first pass used an under-tuned BM25 baseline (nDCG@10 0.228, below the 0.267 literature reference); correcting it — stemming plus title-field weighting, chosen entirely on validation — raises BM25 to 0.315, above the reference, and changes every downstream conclusion. Telemetry available before any action predicts whether unexpanded retrieval is already poor with ROC-AUC 0.649 on frozen test, but the same telemetry vocabulary is markedly worse at the harder, actionable question — will a *specific* expansion help this query — at ROC-AUC only 0.552 ($r = 0.114$, $p < 0.001$). Detection is not actionability, and that gap is this paper’s central finding. Consistent with it, neither static expansion nor a closed-loop agent significantly outperforms corrected BM25 (static: +0.0007, agent: +0.0009 nDCG@10, CIs crossing zero), nor does the agent beat static (+0.0002); reranking alone, with no expansion, significantly *hurts* (−0.0302, $p < 0.001$), so expansion’s role here is largely offsetting harm the reranker introduces. This is not for lack of headroom: an oracle with perfect per-query knowledge gains +0.0246 nDCG@10 over BM25 ($p < 0.001$), of which the verifier captures only 2–13% across intervention budgets. Fault-injection recovery is fault-specific and, via a causal feature-perturbation analysis, mechanistically explained: a grounding-corruption fault is significantly recovered (+0.0099), an entity-injection fault is detected almost every time but flagged almost indiscriminately (net effect not significant), and a reranker-disabling fault is never recovered (recall 0.000) because the fault *increases* the verifier’s healthy-state score rather than merely evading weak features. Generalization is worse still outside the calibration split: action-utility AUC falls to 0.148 (HEAD) and 0.214 (TORSO), both below chance. We report a real implementation defect found mid-study and its correction. Runtime observability reliably detects retrieval failure, but the stronger property autonomous recovery needs — predicting whether a specific intervention will help — is not delivered by the same telemetry, under a protocol designed so that finding is trustworthy.

Index Terms—query expansion, information retrieval, agentic systems, AI observability, retrieval-augmented generation, large language models, site reliability engineering

I. INTRODUCTION

Large language models can generate useful pseudo-documents, alternative phrasings, and disambiguating context for a search query

[1], [2], and corpus-grounded generation — conditioning the rewrite on retrieved evidence rather than the query text alone — improves on naive query-only rewriting. Existing work establishes that LLM query expansion *can* help, not when an operational system should trust an expansion enough to use it. In practice, expansion is applied unconditionally or not at all; neither is safe, since unconditional rewriting is reliably net-harmful on a measurable fraction of queries, particularly underspecified or already-well-served ones.

We treat this as a control problem rather than a generation problem. Borrowing the observe-diagnose-act-recover discipline of site reliability engineering [3] and the trace/metric vocabulary of OpenTelemetry [4], [5], we build an agent whose action space is a small set of interpretable, corpus-grounded expansion operators, whose control loop alternates actions with structured observations as in ReAct [6], and whose accept, roll-back, replan, and abstain decisions are driven by runtime telemetry rather than the language model’s own self-assessment. The central question is not whether an LLM can produce a useful expansion — that is established — but whether *runtime observability* can make expansion selective, recoverable, and cost-aware. Answering it surfaces a distinction the selective query expansion and query performance prediction literature has not, to our knowledge, isolated: predicting that retrieval is failing and predicting that a specific intervention will fix it are different tasks, and telemetry good at the first can be useless at the second. This paper reports what we found, including where the answer was no.

A. Research Questions

RQ1 (Detection vs. Actionability). Can qrel-free retrieval and agent-runtime telemetry (a) identify that unexpanded retrieval is already unhealthy, and (b) predict whether a *specific* expansion action is likely to help or harm — and, if both are possible, is the second any easier than the first once the first is solved?

RQ2 (Control). Does an agent that conditions expansion actions on that telemetry outperform an otherwise identical *static*, failure-conditioned expansion pipeline — the same retriever, grounding mechanism, generator, and fusion method, differing only in whether a closed control loop can observe an action’s consequences and recover from a bad one?

RQ3 (Recovery). After a harmful rewrite or a simulated system-level fault, can observability-driven verification detect the degradation and select an appropriate recovery action to restore retrieval quality, under a bounded action budget — and does this hold even for a verifier whose general-purpose detection power (RQ1) is only weakly informative?

B. Contributions

We contribute (1) an agent architecture treating LLM query expansion as a bounded, typed-tool action inside a closed observability loop, under a frozen train/validation/test protocol (Section IV-B) where every retrieval-config, verifier, and threshold decision is made

on validation before one single-touch test evaluation; (2) a corrected BM25 baseline, selected entirely on validation, moving our TripClick TAIL retriever above the published reference and changing every downstream conclusion; (3) a direct comparison, from one telemetry vocabulary, of failure-*detection* against intervention-*utility* prediction, showing these differ and only the first is reliable; (4) a frozen-test evaluation of the closed-loop agent against a matched static comparator, with an oracle and selectivity-frontier analysis quantifying captured headroom; (5) a fault-injection study with a causal feature-perturbation analysis explaining, not just reporting, which faults a qrel-free verifier recovers, plus a cross-split check showing that ranking degrades below chance outside its calibration split; and (6) an honest account of a real implementation defect, found and corrected mid-study.

II. RELATED WORK

LLM query expansion. Query2doc [1] generates pseudo-documents from an LLM and appends them to the query; MuGI [2] generates and fuses multiple expansions. Both treat expansion as unconditional, one-shot generation, conditioning on neither retrieval-state telemetry nor a harmful rewrite’s recovery.

Agentic control loops. ReAct [6] interleaves reasoning with actions and observations, the template for our observe-act-observe structure, but with a much smaller action space: typed retrieval operators rather than free-form tool calls, keeping the policy interpretable under a non-LLM controller.

Observability and reliability engineering. OpenTelemetry’s trace/metric/log separation and GenAI semantic conventions [4], [5] supply a naming discipline for our telemetry. Google’s SRE framing of an SLO as a target range, not a metric to optimize [3], motivates our Agent Search SLO: maximize quality subject to harm, cost, and latency frozen *before* evaluation.

Query performance prediction and selective expansion. Classic query performance prediction [7] and selective query expansion [8] already ask whether a query is hard, or likely to benefit from expansion, from retrieval signal alone. We split that into two predictive targets sharing one telemetry vocabulary — is retrieval failing, versus will *this* action fix it — and show qrel-free telemetry answers the first far better (Section V).

Benchmarks. TripClick [9] has genuine head/torso/tail query-frequency structure, unlike benchmarks that treat “hard” as a synonym for low-BM25-score, which is why it is our sole evaluation corpus.

III. SYSTEM DESIGN

A. Agent State and Control Loop

The agent’s state at step t is $s_t = (q_0, q_t, R_t, O_t, H_t, B_t)$, where q_0 is the immutable original query, q_t the active search expression, R_t the current ranking, O_t the observability state, H_t the episode’s action/observation history, and B_t the remaining action budget. The control loop is OBSERVE $\rightarrow$ DIAGNOSE $\rightarrow$ ACT $\rightarrow$ RETRIEVE $\rightarrow$ OBSERVE $\rightarrow$ VERIFY $\rightarrow$ {ACCEPT, ROLLBACK, REPLAN, STOP}. The critical property this loop is built to test is that the second decision (accept/rollback/replan/stop) depends on evidence the agent’s own first action created — as opposed to a static pipeline that receives the same query, initial ranking, evidence, generator, and verifier, but executes exactly one predetermined route. Isolating this dependency, not simply comparing against unmodified BM25, is the paper’s central experimental contrast (Section VI).

The action space is deliberately small and interpretable: $A = \{ \text{NOEXPAND, VOCABULARY, CONTEXT, AMBIGUITY, ENTITYNORMALIZE, ROLLBACK, REPLAN, STOP} \}$. VOCABULARY maps informal user language onto corpus-compatible terminology; CONTEXT adds a corpus-supported missing facet to an underspecified query; AMBIGUITY disambiguates toward

the interpretation the evidence best supports; ENTITYNORMALIZE resolves aliases, abbreviations, and canonical-name mismatches. The LLM is invoked through exactly one typed tool, `generate_expansion(action, query, evidence)`; it is never the sole arbiter of its own output’s use — that judgment is a separate, non-generative controller’s (Section III-D).

B. Observability Model

Every query episode produces one trace with spans `retrieve_original` $\rightarrow$ `build_evidence` $\rightarrow$ `plan_action` $\rightarrow$ `generate_expansion` $\rightarrow$ `retrieve_expanded` $\rightarrow$ `rerank` $\rightarrow$ `verify` $\rightarrow$ `recover_or_accept` $\rightarrow$ `fuse`. O_t spans three telemetry families. *Retrieval*: top score, score margin, normalized score entropy, query-term coverage, mean rare-term IDF, entity overlap, top- k ranking overlap against the previous state, reranker/BM25 disagreement, lexical result coherence, and ranking stability under a small controlled query perturbation. *Agent*: selected operator, trajectory length, previous verifier score, rewrite-to-original semantic drift, introduced entities, repeated actions, remaining budget. *System*: per-tool latency, token counts, tool errors/timeouts. Every feature is computable without relevance judgments, which is what makes the verifier usable at inference time.

C. Failure Taxonomy

We label episodes into ten failure categories (vocabulary mismatch, underspecification, ambiguity, entity mismatch, intent drift, confirmation bias, retriever disagreement, tool degradation, budget exhaustion, healthy retrieval), each with an observable telemetry signature and preferred recovery action, motivating Section VII’s fault-injection protocol.

D. Verifier and Controller

The verifier, `verify(previous_state, new_state)`, produces a scalar judgment of whether an action’s consequences look like an improvement, using telemetry alone. Rather than hand-picking its weights, we calibrate them against click-derived qrels (Section IV): for a sample of episodes, we execute one expansion action, compute the true nDCG@10 change against those labels, and fit the verifier’s weights to predict that change from the telemetry features of Section III-B. Click labels are available only during this offline calibration; the deployed verifier never sees one, and is treated as an artifact — trained by a dedicated script, loaded at inference time — so recalibration is a script re-run, not a source edit. The recovery controller is a small deterministic finite-state machine: below a calibrated harmful threshold the agent rolls back to the pre-action ranking; above a calibrated accept threshold it fuses the new ranking via reciprocal-rank fusion; in between it replans with an untried operator, capped at two model calls per query. Both thresholds are calibrated empirically (Section IV), not guessed.

IV. EXPERIMENTAL SETUP

A. Datasets

TripClick [9] is the primary benchmark: 1,523,878 documents, split evenly across head/torso/tail frequency buckets, each further split into train/validation/test (1,175 queries per split per bucket). All results concern the TAIL split unless stated otherwise (Section VIII). Labels are click-derived, not human-assessed (binary raw clicks, matching the label type this split’s published reference point uses [9]); we use “qrel” and “click label” interchangeably, neither implying editorial judgment.




Fig. 1. The observability-guided control loop. Telemetry O_t (orange) drives both operator selection and the accept/rollback/replan/stop decision. The static comparator (Section VI) runs the same chain once, with no path back from VERIFY to the LLM tool.

B. Experimental Protocol

A first pass at this study calibrated the verifier and thresholds on the same episodes later used to report RQ1–RQ3 — an in-sample characterization unsuited to an out-of-sample generalization claim. We froze a stricter protocol instead, in a committed config rather than left implicit: TAIL *validation* may be used for any retrieval-config, verifier, or threshold decision; TAIL *test* may be touched at most once per research question, purely for evaluation, after every upstream decision is frozen, with no decision afterward conditioned on a test-set outcome. This discipline directly produced the paper’s first finding, below.

C. A Corrected Retrieval Baseline

BM25 over an unstemmed FTS5 index scored nDCG@10 0.228 on TAIL test — 15% below the 0.267 reference for this exact split and label type [10]. Selecting entirely on validation, we found two config choices that close this gap: Porter stemming (val nDCG@10 0.235 $\rightarrow$ 0.270) and weighting the title field ten times body weight in a title/body-split FTS5 index (val nDCG@10 climbing to a clear plateau at that ratio, 0.327, MRR peaking there and declining beyond it). A single, one-time sanity check of the frozen config against withheld test — not a selection criterion — gives nDCG@10 0.315, *above* the published reference. We report this because an under-tuned baseline mechanically inflates the apparent value of anything compared against it: every downstream conclusion in this paper changed once this fix was in place, a precondition for the rest of the study rather than an implementation footnote.

D. Implementation

Sparse retrieval uses BM25 over a SQLite FTS5 index [11] rather than a JVM Lucene index, forced by a JVM crash under Pyserini [12] on our target hardware; two further silent-corruption risks surfaced at this corpus’s scale, in the same category as Section IV-C’s baseline gap and Section VII-A’s telemetry defect: FTS5’s default MATCH combining operator is *conjunctive* (collapsing recall without an error), and FTS5 indexes only the MATCH column (an unindexed companion lookup degrades to a full table scan, an $18\times$ slowdown at 1.5M rows). Cross-encoder reranking (cross-encoder/ms-marco-MiniLM-L6-v2 [13]) is applied to the top 100 BM25 candidates, reranking the top 50. Generation uses a locally hosted, 4-bit quantized 2–3B-parameter model with a hard cap of two calls per query episode and a 32–64 token structured output, so the language model is never the entire decision surface, only one typed tool within it.

TABLE I

MODEL SEARCH ON TAIL *validation* ($n = 1147\text{--}1168$): PEARSON r VS. TRUE NDCG@10 CHANGE. SELECTS A MODEL, NOT THE GENERALIZATION RESULT (TABLE II).

Model, feature family	Val Pearson r
v0 hand-picked weights (6 features)	+0.020 (n.s., $p=0.50$)
Linear, 10 features (lexical+reranker)	−0.014
GBDT, 10 features	+0.001
Linear, 13 features (+dense embedding)	−0.015

E. Verifier Calibration Methodology

Following Section IV-B, we collect calibration episodes (one expansion action per query, true nDCG@10 change against click labels) on TAIL *validation* ($n = 1147\text{--}1168$ depending on feature family), and compare three feature families — (i) six retrieval-telemetry deltas; (ii) the same six plus four reranker- and length-derived features; (iii) family (ii) plus three dense-embedding features — against three model classes each (fixed-weight baseline, 5-fold CV linear, 5-fold CV GBDT), under Pearson correlation (Table I), to characterize the detection ceiling rather than accept the first plausible fit. This selects a model and is never itself reported as evidence the signal generalizes (Section V reports that separately, from one frozen touch of test). Recovery thresholds are calibrated the same way, by checking which candidate cutoffs on validation separate true-harmed episodes from the rest.

F. Metrics

Retrieval quality is nDCG@10, Recall@100, and MRR, plus the fraction of queries helped/unchanged/harmed relative to a baseline. For the recovery study (Section VII) we report Detection Recall $P(\text{intervened} \mid \text{harmful})$, False Alarm Rate $P(\text{intervened} \mid \text{healthy})$, and recovery gain in mean nDCG@10 with vs. without recovery on identically corrupted episodes. We additionally report a retrospective *oracle ceiling*: per query, $\max(\text{baseline}, \text{treatment})$ under true click labels, diagnosing whether headroom exists independent of whether our qrel-free verifier can find it (Section VI). Significance is via paired bootstrap (10,000 resamples) throughout.

V. RQ1: DETECTION

Nothing clears noise on validation: every candidate — a hand-picked baseline, linear and gradient-boosted fits on the deployed 10-feature set, and a fit augmenting it with three dense-embedding features — lands at or below $|r| = 0.02$. We select the 10-feature linear model anyway, the only candidate fit on the corrected pipeline’s

TABLE II

RQ1'S FROZEN RESULTS: BOTH MODELS EVALUATED *once* ON WITHHELD TAIL TEST ($n = 1157$), NEVER TOUCHED DURING SELECTION.

Metric	Test value
Action utility: Pearson r vs. true Δ nDCG@10	0.114 ($p < 0.001$)
Action utility: ROC-AUC (detecting harm)	0.552
Action utility: average precision (prev. 0.330)	0.364
Failure detection: ROC-AUC (val CV 0.618)	0.649

TABLE III

RQ2, FROZEN RESULT, FULL TRIPCLICK TAIL TEST ($n = 1175$), CORRECTED BASELINE. STATIC QE AND AGENT SHARE RETRIEVER, GROUNDING, GENERATOR, VERIFIER, FUSION; AGENT ADDS A CLOSED OBSERVE-VERIFY-RECOVER LOOP.

System	nDCG@10	Recall@100	MRR
BM25 only	0.3152	0.7750	0.3014
BM25 + rerank, no QE	0.2850 [†]	—	—
Static QE	0.3159	0.7759	0.2993
Agent	0.3161	0.7770	0.3015

[†]: significantly below BM25 only, -0.0302 [$-0.0447, -0.0154$], $p < 0.001$.

feature distribution rather than a stale one, and deploy it *despite* its weak validation estimate, per Section IV-B: no decision here is conditioned on test performance.

The frozen action-utility model generalizes better than its noisy validation estimate suggested: $r = 0.114$ on test is small but real ($p < 0.001$, $R^2 \approx 1.3\%$), modestly above no-skill baselines (Figure 2). We read the val/test gap honestly rather than as vindication: cross-validation on $\sim 1,150$ episodes is a noisy estimator for an effect this small, and one larger test set landing positive does not retroactively make the selection sound.

A. Detection Is Not Actionability

That action-utility model answers a hard question: given telemetry before *and* after a specific action, did it help? A logically prior, cheaper question uses only pre-action telemetry: is unexpanded retrieval already unhealthy? We build a second model from the same eight retrieval-telemetry features, labeling a query “poor” when BM25 surfaces zero relevant documents in the top 10 (test prevalence 0.450), selected by 5-fold CV on validation and touched once on frozen test, identically to Section IV-E’s protocol — Table II’s last row.

The gap is large and consistent across val and test: telemetry genuinely informative about *whether* retrieval is failing (AUC = 0.649) is markedly weaker at predicting whether a *specific candidate action* fixes it (AUC = 0.552) — both frozen single-touch test numbers from the same features and protocol, not the same result at different noise levels. **RQ1’s answer:** qrel-free telemetry reliably detects that something is wrong, but is only weakly informative about whether a proposed fix will work — the gap RQ2’s control result and RQ3’s fault-specific pattern trace back to.

VI. RQ2: CONTROL

Once BM25 itself is corrected (Section IV-C), the picture changes entirely: neither Static QE nor Agent significantly outperforms BM25 alone (static: $+0.0007$ nDCG@10, 95% CI [$-0.0066, 0.0079$]; agent: $+0.0009$, 95% CI [$-0.0069, 0.0087$]), and Agent does not significantly outperform Static QE ($+0.0002$, 95% CI [$-0.0027, 0.0032$]; changes 3.9% of outcomes, 2.1% helped, 1.7% harmed). The agent still costs 24% more model calls (1.22 vs. 0.98 per query) and 36% higher mean latency (15.1 s vs. 11.1 s, excluding sleep-inflated episodes) for no offsetting gain. Isolating the reranker alone — corrected BM25 candidates, cross-encoder reranking, no expansion — shows reranking by itself is significantly *harmful* here

TABLE IV

ORACLE SELECTIVITY CEILING, SAME EPISODES AS TABLE III. ORACLE PICKS $\max(\text{BM25}, \text{TREATMENT})$ PER QUERY, TRUE CLICK LABELS.

Comparison	nDCG@10 gain	95% CI
Oracle-static vs. BM25	$+0.0246$	[0.0197, 0.0296]*
Oracle-static vs. Static QE	$+0.0239$	[0.0191, 0.0292]*
Oracle-agent vs. Agent	$+0.0276$	[0.0226, 0.0332]*

*: significant, $p < 0.001$.

TABLE V

SELECTIVITY FRONTIER, TAIL TEST: MEAN NDCG@10 GAIN OVER ALWAYS-BM25 WHEN THE TOP $X\%$ OF QUERIES BY POLICY RANK GET TREATMENT, REST AT BM25. % ORACLE: VERIFIER GAIN AS A FRACTION OF ORACLE GAIN.

Budget	Verifier gain	Oracle gain	% Oracle
5%	0.0005	0.0186	2.4%
10%	0.0007	0.0245	2.7%
20%	0.0023	0.0250	9.1%
40%	0.0033	0.0250	13.1%
100%	0.0007	0.0007	100%*

*At 100% budget every policy treats every query, so ranking cannot matter; verifier and oracle necessarily coincide with Table III’s unconditional Static QE result.

(-0.0302 , $p < 0.001$): the cross-encoder, trained on general web relevance, misorders health-domain results. Static QE and Agent both rerank *after* expanding, yet land statistically even with unranked BM25, so expansion’s practical role here is largely canceling harm the reranker introduces, not adding quality on top of a neutral reranking step. **On natural queries, neither expansion flavor moves the needle once the retriever is competitive** — materially different from what a weak baseline would have supported.

A. Is There Any Headroom at All?

A null result here is ambiguous: either no meaningful control opportunity exists, or one exists but this verifier cannot find it. We resolve this with a retrospective oracle (Section IV): per query, take whichever of BM25 or the treatment ranking true click labels say is better — not deployable, since it uses labels unavailable at inference time, but useful to bound the second explanation.

The answer is unambiguous: substantial headroom exists. A perfect selective controller would gain $+0.0246$ nDCG@10 over BM25 — roughly $35\times$ what the deployed verifier’s Static QE actually captures ($+0.0007$) — both significant at $p < 0.001$. This is not a corpus with nothing to find: static expansion helps 11.6% of queries and harms 10.6% relative to BM25 (Table III’s per-query comparison) — a real, roughly-balanced mix of genuine gains and harms a perfect selector could exploit. The failure is squarely in *detection*: RQ1’s verifier, at $r = 0.114$ on frozen test, is far too weak to discriminate the helped 11.6% from the harmed 10.6% well enough to move aggregate quality — converting RQ2’s null result from “maybe there’s nothing here” into “there is something here, and our qrel-free detector cannot see it.”

At every budget below universal treatment, the verifier captures only a small fraction of available headroom — 2.4% at 5%, rising to 13.1% at 40%, never close to the oracle. Random selection at these budgets is statistically indistinguishable from zero gain (not shown), so the ranking is not worthless, but it is far short of what perfect knowledge extracts.

Table VI shows a controller that intervenes often (38.9% ROLL-BACK or REPLAN) without those interventions paying off in aggregate: the frozen verifier’s thresholds (Section IV-E) are wide, since nothing on validation cleanly separated harmful from healthy episodes, so the controller second-guesses frequently but not accurately. **RQ2’s answer is no**, with the oracle analysis clarifying




Fig. 2. RQ1’s frozen result: ROC and precision-recall curves for the selected verifier’s score against true harm (nDCG@10 $\Delta < 0$) on withheld TAIL test ($n = 1157$, single touch). The curve sits consistently, but only modestly, above the no-skill baseline (dashed) in both panels.

TABLE VI
AGENT CONTROLLER-DECISION FREQUENCY, FULL TRIPCLICK TAIL TEST SET ($n = 1175$).

Decision	Count	%
ACCEPT	699	59.5
ROLLBACK	331	28.2
REPLAN	127	10.8
ABSTAIN (no expansion)	18	1.5

why: an observability-guided closed loop built on this verifier does not outperform a matched static pipeline, not because no control opportunity exists, but because this verifier cannot exploit it.

VII. RQ3: WHICH FAILURES ARE RECOVERABLE?

RQ1 and RQ2 leave open whether observability has *any* practical value, or whether a weak general detector is simply useless. We test a narrower claim: that the same weak verifier can still serve as a safety net against *large, obvious* failures even though it cannot reliably judge *subtle, natural* ones. We inject three fault types at their natural point of corruption — an unsupported entity appended to the expansion, a grounding passage swapped for a non-supporting one, and the reranker disabled — and compare, on identical corrupted episodes, an always-accept ablation against the real calibrated controller.

A. A Defect That Changed the Answer

An early pass at this study found the reranker-disabled fault essentially undetectable (recall 0.185) while the other two looked strong. Tracing this revealed a real defect: the reranked ranking was computed, at real latency cost, and fed into telemetry, but never actually applied to the ranking the agent accepted or evaluated — disabling the reranker was mechanically a no-op regardless of verifier quality. Fixing it exposed a second, entangled defect: the verifier’s telemetry was still computed from the pre-rerank ranking, judging a different top- k than the one being accepted. We corrected both and reran calibration to confirm the deployed verifier did not need retraining on the corrected telemetry (corrected-pipeline Pearson $r = 0.146$ vs. the original 0.134). We report this in full because a component computed purely for observability is architecturally easy to leave disconnected from the outcome it explains, and the resulting error looked coherent — unrealistically strong — until an implausible near-zero effect prompted investigation.

B. Frozen Results

The three fault types split into three distinct outcomes, each with a different explanation, so “is recovery useful” has no single answer — only a per-failure-type one. **Entity injection** is detected almost every time (recall 0.974) but flagged almost indiscriminately (false alarm 0.939): the controller treats nearly any entity-injected rewrite as suspect whether or not the injection was actually benign, and rolling back that many healthy episodes offsets the genuine catches, leaving a positive but not significant net effect. **Grounding corruption** is the one clean recovery success (+0.0099, $p < 0.05$): moderate

TABLE VII
RQ3, FROZEN RESULT: FAULT INJECTION, CORRECTED PIPELINE, TAIL TEST ($n = 1157$; 18/1175 SKIPPED, NO EXPANSION). GAIN IS MEAN NDCG@10 WITH MINUS WITHOUT RECOVERY ON IDENTICALLY CORRUPTED EPISODES; CI FROM 10,000-RESAMPLE BOOTSTRAP. DOMINANT CONTRIBUTOR: FEATURE j WITH LARGEST $|w_j \Delta x_j|$ (BELOW).

Fault type	Rec.	FA	Gain (95% CI)	Dominant contrib.
Entity injection	0.974	0.939	+0.0081 [−0.003, +0.019]	coverage_gain −0.054
Grounding corruption	0.528	0.469	+0.0099 [+0.002, +0.018]*	topk_overlap +0.002
Reranker disabled	0.000	0.002	+0.0000	disagreement +0.126

CIs straddle zero except *: significant, positive. FA: false alarm.

recall (0.528) paired with a comparable false alarm rate (0.469) is discriminating enough, on this fault type specifically, to produce a real net gain. **Reranker disabled** is never detected at all (recall 0.000, false alarm 0.002): the controller essentially always accepts.

Why, causally. Fitted weights alone (Section IV-E) show which features the verifier ignores, not what a fault *does* to its score. We measure this directly: for each fault, $\Delta x_j = \mathbb{E}[x_j^{\text{fault}} - x_j^{\text{clean}}]$ per feature j , and $\text{Contribution}_j = w_j \cdot \Delta x_j$ using the frozen verifier’s weights ($n=79/\text{fault}$), sign relative to its healthy-state direction (Table VII, last column). Entity injection’s dominant feature pushes the score sharply toward “unhealthy,” consistent with its high recall; the blunt false-alarm rate is a threshold problem, not feature-blindness. Grounding corruption’s dominant contribution is an order of magnitude smaller, matching its near-coin-flip recall. Reranker disabled is mechanistically distinct: its dominant contribution is large and *positive*, from disagreement moving in the direction the verifier reads as improvement — the fault does not merely evade weak reranker-derived features (reranker_top_score +0.0007), it makes the corrupted state score higher than the clean one, so the controller is actively misled, not merely blind.

RQ3’s answer is that recoverability is fault-specific, and the attribution shows why: content-level corruption is recoverable when it perturbs a strongly-weighted feature the unhealthy way (grounding), a wash when detection is blunt (entity injection), and a system-level fault is disguised as an improvement when it moves a strongly-weighted feature the *wrong* way (reranker) — a verifier is not uniformly weak or strong across failure types, it can be actively adversarial to a fault it was never calibrated against.

VIII. DOES THE VERIFIER GENERALIZE BEYOND TAIL?

Every calibration decision so far used TAIL validation only (Section IV-B); we did not test whether the frozen verifier’s action-utility ranking holds on TripClick’s HEAD and TORSO frequency buckets, which it never saw during selection. Using the identical frozen retrieval config and verifier, no recalibration, we run BM25 and Static QE on each bucket’s full test split.

Static QE stays statistically indistinguishable from BM25 on every bucket, matching TAIL, and oracle headroom exists on every bucket too, comparable in magnitude to TAIL’s. What does not transfer is the verifier itself: action-utility AUC falls from 0.552 on TAIL to 0.214 on TORSO and 0.148 on HEAD — both *below* chance,

TABLE VIII
CROSS-SPLIT ROBUSTNESS, FULL TEST SPLITS ($n = 1175$ EACH), FROZEN
VERIFIER AND THRESHOLDS UNCHANGED.

Split	BM25	Static QE	Oracle gain	AUC
HEAD	0.2282	0.2293 (NS)	+0.0153*	0.148
TORSO	0.2485	0.2506 (NS)	+0.0187*	0.214
TAIL	0.3152	0.3159 (NS)	+0.0246*	0.552

NS: static-vs-BM25 CI crosses zero. *: oracle-vs-static significant,
 $p < 0.001$. AUC is the frozen verifier’s action-utility AUC on that split.

so the verifier’s ranking is not merely uninformative outside TAIL but anti-correlated with true utility there. A verifier calibrated on one frequency regime does not just fail to add value in another; its ranking inverts, which a deployment relying on this signal outside its calibration split would not detect without exactly this kind of check.

IX. DISCUSSION

Read together, RQ1–RQ3, the oracle analysis, and the cross-split check sketch one finding, not five independent ones. Section V-A is load-bearing: the same telemetry vocabulary detects that retrieval is failing reasonably well (AUC 0.649) but is only weakly informative about whether a specific action will fix it (AUC 0.552). RQ2 shows the practical consequence — a closed control loop built on the weaker signal adds nothing over a matched static pipeline — and the oracle and selectivity-frontier analyses (Section VI-A) rule out “no opportunity exists” as the explanation: real headroom is available (+0.0246 nDCG@10 to a perfect selector) and the verifier reaches at most 13% of it at any tested budget. RQ3’s causal attribution shows this is not one uniform weakness but a per-fault property, sharp enough in one case (reranker disabled) that the verifier actively reads the fault as an improvement, because it moves a strongly-weighted feature the wrong way. Section VIII shows the same fragility along a different axis: the verifier’s ranking, held fixed, goes from weakly informative on TAIL to below-chance on HEAD and TORSO, so whatever this telemetry captures is tied to the regime it was calibrated on.

Two things here were found only by refusing a convenient result: the BM25 baseline (Section IV-C) looked adequate until checked against a published reference, reversing RQ2’s conclusion; the reranker-wiring defect (Section VII-A) looked like a strong recovery result until one fault type’s implausible near-zero effect prompted investigation. Both are cases where a component is architecturally easy to leave unchecked against an external reference, and the resulting error looks coherent rather than broken. A frozen protocol with a mandatory sanity check (Section IV-B) defends against the first; the second we only found by distrusting a result that seemed too good, which no protocol automates.

X. LIMITATIONS

The oracle ceiling (Section VI-A) is retrospective and uses test-set click labels directly — a valid upper-bound diagnostic, never a deployable method. The generator is a locally hosted, quantized 2–3B parameter model under a strict two-call budget; a larger generator might change RQ1’s ceiling or RQ2’s null result. Cross-split robustness (Section VIII) covers TripClick’s three buckets only; other benchmarks remain untested, so the AUC collapse could be frequency-specific or reflect broader corpus-shift sensitivity. Click-derived qrels are an imperfect relevance proxy inherited from the benchmark. The corrected baseline’s title-weight ratio ($10\times$ body) came from a genuine validation-set plateau, not cherry-picking, but may not transfer beyond TripClick’s title-informative abstracts. The causal attribution in Section VII uses the frozen verifier’s linear weights as a local description of its own behavior. A prospective check on one held-out fault (query-term removal, predicted before running from the weights alone; `scripts/12`) got the qualitative direction right — detectable, entity-injection-like profile — but not

the specific dominant feature: the mechanism generalizes qualitatively, not as a precise single-feature account. The failure-detection model (Section V-A) uses one binary definition of “poor retrieval”; a graded definition could shift its measured AUC.

XI. CONCLUSION

We evaluated a closed-loop, observability-guided query-expansion agent under a corrected protocol separating model selection from final evaluation — surfacing the paper’s first finding: an honestly-selected BM25 fix inverted every downstream conclusion. Runtime observability reveals retrieval failures reliably (detection AUC 0.649), but autonomous recovery needs a stronger property the same telemetry does not deliver: predicting whether a specific intervention will help (action-utility AUC 0.552, below chance outside its calibration split). That gap, not any single agent-vs-static comparison, is why the closed-loop agent does not outperform static expansion and neither outperforms BM25, despite real headroom the oracle and selectivity-frontier analyses show going uncaptured. The recovery study shows this is a per-fault, mechanistically explicable property — in one case the verifier is not merely blind to a fault but actively misled by it. We report both implementation defects found by refusing a convenient result, in the belief that the discipline to catch them is worth more here than a flattering number.

ACKNOWLEDGMENT

The authors thank the TripClick maintainers for access to the dataset under its non-commercial research license.

REFERENCES

- [1] L. Wang, N. Yang, and F. Wei, “Query2doc: Query expansion with large language models,” in *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, 2023. [Online]. Available: <https://aclanthology.org/2023.emnlp-main.585/>
- [2] L. Zhang, Y. Wu, Q. Yang, J. Lin *et al.*, “Exploring the best practices of query expansion with large language models,” in *Findings of the Association for Computational Linguistics: EMNLP 2024*, 2024. [Online]. Available: <https://aclanthology.org/2024.findings-emnlp.103/>
- [3] Google, “The site reliability workbook: Defining SLO,” 2026. [Online]. Available: <https://sre.google/sre-book/service-level-objectives/>
- [4] OpenTelemetry Authors, “OpenTelemetry signals,” 2026. [Online]. Available: <https://opentelemetry.io/docs/concepts/signals/>
- [5] —, “OpenTelemetry GenAI semantic conventions 1.44.0,” 2026. [Online]. Available: <https://opentelemetry.io/docs/specs/semconv/>
- [6] S. Yao, J. Zhao, D. Yu, N. Du, I. Shafraan, K. Narasimhan, and Y. Cao, “React: Synergizing reasoning and acting in language models,” *arXiv preprint arXiv:2210.03629*, 2022. [Online]. Available: <https://arxiv.org/abs/2210.03629>
- [7] S. Cronen-Townsend, Y. Zhou, and W. B. Croft, “Predicting query performance,” in *Proceedings of the 25th Annual International ACM SIGIR Conference on Research and Development in Information Retrieval*, 2002. [Online]. Available: <https://dl.acm.org/doi/10.1145/564376.564429>
- [8] —, “A framework for selective query expansion,” in *Proceedings of the Thirteenth ACM International Conference on Information and Knowledge Management*, 2004. [Online]. Available: <https://dl.acm.org/doi/10.1145/1031171.1031220>
- [9] N. Rekabsaz, O. Lesota, M. Schedl, J. Brassey, and C. Eickhoff, “TripClick: The log files of a large health web search engine,” 2021. [Online]. Available: <https://tripdatabase.github.io/tripclick/>
- [10] —, “Establishing strong baselines for TripClick health retrieval,” in *European Conference on Information Retrieval*, 2022. [Online]. Available: <https://arxiv.org/abs/2201.00365>
- [11] SQLite Consortium, “SQLite FTS5 extension,” 2026. [Online]. Available: <https://www.sqlite.org/fts5.html>
- [12] J. Lin, X. Ma, S.-C. Lin, J.-H. Yang, R. Pradeep, and R. Nogueira, “Pyserini: Reproducible IR research with sparse and dense representations,” 2026. [Online]. Available: <https://github.com/castorini/pyserini>
- [13] Hugging Face, “cross-encoder/ms-marco-MiniLM-L6-v2,” 2026. [Online]. Available: <https://huggingface.co/cross-encoder/ms-marco-MiniLM-L6-v2>